


Introduction to Python Programming (3/3)




Institute of Medical Informatics
University of Lübeck

Last session

- Introduction of basic Python concepts
- Basic Python exercises
- Implementation of a Tic Tac Toe game

```
1 import pygame
2 from pygame.locals import *
3 import random
4
5 # Class which describes the board of TicTacToe
6 class Board:
7     # list with all elements of the board
8     list = None
9
10    def __init__(self):
11        self.empty_board()
12
13    # function to empty all values from the board
14    def empty_board(self):
15        # initialize a list consisting of 3 lists with 3 '0's
16        self.list = [[0]*3 for i in range(3)]
17
18    # function to make a move. Value of player is set to position i,j of the board.
19    def set_token(self, player, i, j):
20        if (self.list[i][j] == 0):
21            self.list[i][j] = player
22            return True
23        return False
24
25    # function to make a random move by the computer.
26    def computer_random_move(self, player = 2):
27        found_free_place = False
28
29        while (not found_free_place):
30            i = random.randint(0,2)
31            j = random.randint(0,2)
32
33            if (self.list[i][j] == 0):
34                self.list[i][j] = player
35                found_free_place = True
```



Contents of the session

- Array and list indexing
- **[Exercise 1]** Shuffling in unison
- Multidimensional array manipulation
- **[Exercise 2]** Standard normalization
- Evaluation of supervised learning algorithms
- **[Exercise 3]** Confusion matrix, precision and recall
- k-means clustering
- **[Exercise 4]** k-means implementation

Array and list indexing (1/4)

The elements of arrays and lists can be accessed using the **brackets** [...] notation.

Examples:

```
>>> import numpy as np
>>> T = np.arange(5) # Creates array([0, 1, 2, 3, 4])
>>> T[0] # First element of T
0
>>> T[2] # 3rd element of T
2
>>> T[-1] # Last element of T
4
>>> T[-2] # Last-but-one element of T
3
```

Array and list indexing (2/4)

It is possible to access specific elements of an array using **slice indexing**. Given an array or list `T`:

`T[start:end:increment]`

returns all elements of `T` between the indices `start` (included) and `end` (excluded) by intervals of `increment`.

Note 1: `start`, `end` and `increment` are all optional (integer) arguments. If not specified, `increment` is equal to 1 by default.

Note 2: this syntax still works if `T` is a multidimensional array

Array and list indexing (3/4)

Slice indexing examples:

```
>>> import numpy as np
>>> T = np.arange(5) # Creates array([0, 1, 2, 3, 4])
>>> T[1:4] # Elements of T between the 2nd and 4th
array([1,2,3])
>>> T[:3] # All elements of T until the 3rd one
array([0,1,2])
>>> T[2:] # All elements of T from the 3rd to the end
array([2,3,4])
>>> T[1:4:2] # All second elements of T between the 2nd and 4th
array([1,3])
>>> T[:] # Returns all elements of T; equivalent to T[:, :] and T
array([0,1,2,3,4])
```

Array and list indexing (4/4)

Extracting a sub-array can also be performed using **list indexing**. The list can contain either boolean or integers.

Examples:

```
>>> import numpy as np
>>> T = np.array([12,5,-3,7,24])
>>> b = [True,False,False,True,True] # List of booleans
>>> T[b]
array([12,7,24])
>>> idx1 = [3,2,0,4,1] # List of integers with same length than T
>>> T[idx1]
array([7,-3,12,24,5])
>>> idx2 = [2,3,0] # List of integers shorter than the length of T
>>> T[idx2]
array([-3,7,12])
```

Exercise 1: Shuffling in Unison

Write a Python function `shuffleInUnison` which shuffles (i.e. re-orders) the elements of two arrays of same length using the **same** random permutation:

```
shuffledT1, shuffledT2 = shuffleInUnison(T1,T2)
```

With:

- **[input]** T1 and T2: arrays assumed to have same length
- **[output]** shuffledT1 and shuffledT2: randomly shuffled input arrays/lists using the same permutation

Tip: for random related functions, use the `numpy.random` package

Multidimensional array manipulation (1/3)

Some useful options to initialize a multidimensional array:

- Direct initialization with `numpy.array`

```
>>> import numpy as np
>>> T = np.array([[1,2],[3,4]]) # 2x2 array containing the
    values 1, 2, 3, 4
```

- Initialization with numpy functions

```
>>> import numpy as np
>>> T0 = np.zeros((2,2),dtype=int) # 2x2 array to zero
    with integer type
>>> T1 = np.ones((100,50,200),dtype=float) # 3D array to
    one with float type
>>> TRand = np.random.rand(10,20) # 2x2 array of random
    float (default) values
```

Multidimensional array manipulation (2/3)

Multidimensional array **indexing** and **slicing** works the same way than for the 1D case. In particular, for a N-dimensional array T it is possible to use the following syntax to obtain a specific slice of T:

`T[start1:end1:incr1, start2:end2:incr2, ..., startN:endN:incrN]`

Examples:

```
>>> import numpy as np
>>> T = np.random.rand(10,5,20,10,dtype=float) # 4D array of
random float values
>>> s1 = T[-2,3,15,0] # Float element at position (8,3,15,0)
>>> s2 = T[:,3,:,:] # 3D slice of shape (10,20,10)
>>> s3 = T[5,0:5:2,9,:] # 2D slice of shape (3,10)
>>> s4 = T[1:4,::,-1] # 3D slice of shape (3,5,20)
```

Multidimensional array manipulation (3/3)

Numpy arrays are **mutable**: it is possible to change their values after initialization.

Examples:

```
>>> import numpy as np
>>> T = np.zeros((2,2)) # 2x2 array of random float values
>>> T[0,0] = 1; print(T)
array([[1., 0.],
       [0., 0.]])
>>> T[:,1] += 3; print(T) # Increment by 3 the 2nd column of T
array([[1., 3.],
       [0., 3.]])
>>> T[1,:] = -2; print(T) # Set the 2nd line of T to -2
array([[1., 3.],
       [-2., -2.]])
```

Exercise 2: Standard Normalization (1/2)


Multimodal time-series data records can be represented as a 2D array of size $T \times S$ with:

- T : number of timestamps (i.e. duration of the data record)
- S : number of sensor channels

It is usually needed to perform pre-processing operations on the raw data records. **Standard normalization** is one of them:

$$X \leftarrow \frac{X - \mu}{\sigma}$$

with $\mu = \text{mean}(X)$ and $\sigma = \text{std}(X)$



Exercise 2: Standard Normalisation (2/2)

- 1) Initialize a 2D array of random values between -1 and 1 - called `data` - of size $(T, S) = (1000, 3)$
- 2) Multiply each column of `data` with the following respective coefficients: $[10, 0.5, 100]$
- 3) Write a function `standardNormalization` which takes a 2D array as input and returns the input 2D array with a standard normalization **independently** performed on all its sensor channels:

```
normalisedData = standardNormalisation(data)
```

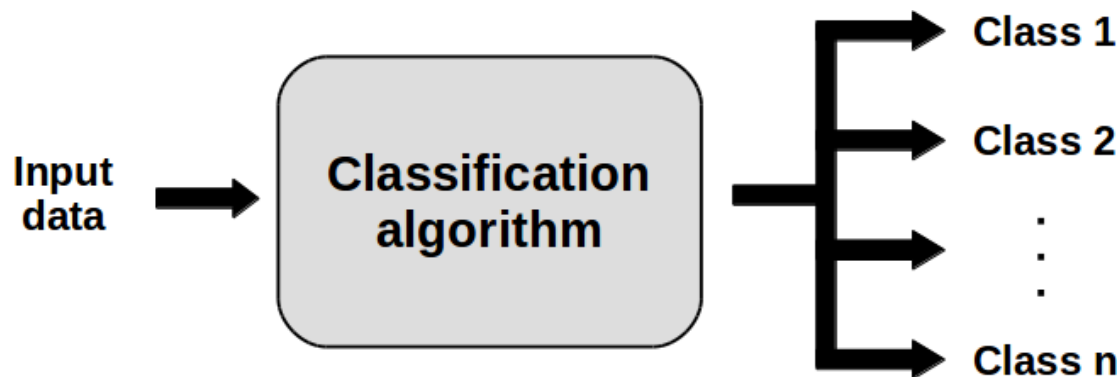
With:

- **[input]** `data`: 2D array (of implied dimensions `nbTimestamps` x `nbSensors`)
- **[output]** `normalisedData`: data with all its sensor channels standard-normalised independently

Test your function on the 2D array `data` previously defined.

Evaluation of supervised learning algorithms (1/5)

A pattern recognition problem can be translated into a **classification** problem, with **classes** = objects to recognize.






The most popular way to obtain a proper classification algorithm is to use a **supervised learning approach**.

Supervised approaches require a **training set** with its associated **labels**.

Evaluation of supervised learning algorithms (2/5)

Each element of a **labeled dataset** is associated to one **label**. A label indicates **to which class the example belongs to**.

<u>Training examples</u>	<u>Associated labels</u>
	butterfly
	dolphin
	leopard
...	...

Integers between 0 and N-1 are usually used as labels, where N is the total number of classes.

For each example, the classifier outputs a **prediction** which is also an integer between 0 and N-1.

Evaluating the classifier can be done by **comparing the "true" labels to the estimated ones**.

Evaluation of supervised learning algorithms (3/5)

One handy tool for the evaluation of classification algorithms:
confusion matrix.

A confusion matrix **counts and compares** the estimations of a trained classification algorithm to the true labels.

		Estimated Labels			
		Class 1	Class 2	...	Class n
True Labels	Class 1	793	34	...	29
	Class 2	105	567	...	78

	Class n	87	14	...	1025

Evaluation of supervised learning algorithms (4/5)

It is possible to compute evaluation metrics from a confusion matrix. For a **binary** classification problem:

- **Accuracy** = $(TP+TN) / (TP+TN+FP+FN)$
- **Precision** = $p = TP / (TP+FP)$
- **Recall** = $r = TP / (TP+FN)$
- **F1-score** = $2*p*r / (p+r)$

True Labels	Estimated Labels	
	Class +	Class -
Class +	TP	FN
Class -	FP	TN

TP = True Positive
FN = False Negative
TN = True Negative
FP = False Positive

Evaluation of supervised learning algorithms (5/5)

For a **multiclass** classification problem (n classes with $n \geq 3$), the **overall precision (recall)** can be obtained by averaging class precisions (respectively class recalls) over the n classes.

The **accuracy** is provided by $\text{Trace}(\mathbf{M}) / \text{Sum}(\mathbf{M})$ with \mathbf{M} confusion matrix.

$\mathbf{M} =$

Estimated Labels					True Labels
	...	Class i	...	Class n	
...	...	FP_i	
Class i	FN_i	TP_i	FN_i	FN_i	
...	...	FP_i	
Class n	...	FP_i	

Exercise 3: Confusion matrix, precision, recall (1/2)

Write a function `confusionMatrix` which computes and returns a confusion matrix given two vectors of true and estimated labels:

```
confMat = confusionMatrix(nbClasses,trueLabels,estimatedLabels)
```

With:

- **[input]** `nbClasses`: integer giving the total number of classes
- **[input]** `trueLabels`: vector of integers containing true labels (elements assumed to be between 0 and `nbClasses-1`)
- **[input]** `estimatedLabels`: vector of integers containing estimated labels (assumed to have the same length as `trueLabels`; elements between 0 and `nbClasses-1`)
- **[output]** `confMat`: 2D array of size `nbClasses x nbClasses` containing the confusion matrix of the classification problem

Exercise 3: Confusion matrix, precision, recall (2/2)

Write the functions `precision` and `recall` which compute and returns the overall precision and recall given a confusion matrix:


```
p = precision(confusionMatrix)
r = recall(confusionMatrix)
```

With:

- **[input]** `confMat`: 2D array of size `nbClasses x nbClasses` containing the confusion matrix of the classification problem
- **[output]** `p / r`: overall precision / recall of the classifier



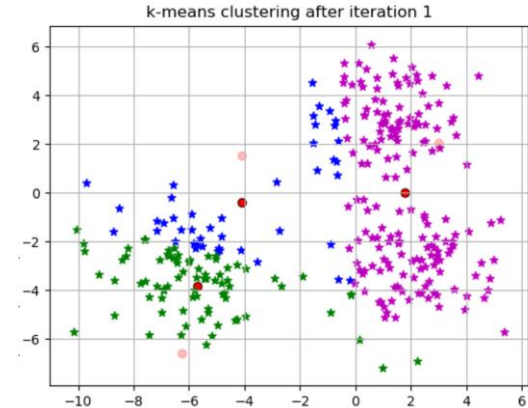
K-means clustering (1/3)

- **Cluster analysis** = one branch of unsupervised learning
- **Goal:** find patterns in data without any external information (labels) to determine the structure of the data
- One of the most famous (and simple) clustering approach is **k-means clustering**.




K-means clustering (2/3)

- **Input:**
 - X: d-dimensional dataset
- **Hyper-parameters:**
 - k: number of clusters to find
 - iterMax: number of maximum iterations
- **Pseudo-code:**
 - Initialise k cluster centers at random
 - While no convergence and $nbIter < iterMax$:
 - Attribute each point of X to its closest cluster center
 - Compute the new centers as the average of the points in the cluster
- **Note:** convergence = the positions of the cluster centers no longer change




K-means clustering (3/3)

k-means clustering after convergence



k-means clustering after iteration 20



- It is possible to mathematically prove that the k-means algorithm **always converges** ...
- ... but not always towards an optimal point!
- The initialisation of the starting cluster centers has a big impact on the **final cluster centers obtained after convergence + convergence time**.
- To make the convergence easier, possible to use **simple heuristic rules** (e.g. boundary on the min/max values for the initial coordinates of the centers)

Exercise 4: K-means implementation

- A main function implementing the algorithm has already been written in `1.3-kMeans-exercise.py`
- Write the code for the following functions:
 - `initialiseCenters(...)`: create k cluster centers at random in a d -dimensional space.
 - `checkConvergence(...)`: check if the k-means algorithm has converged.
 - `attributeCluster(...)`: associate each point of a dataset to its closest cluster center.
 - `computeNewCenters(...)`: compute new cluster centers given a dataset and cluster labels.
- All details regarding expected input and output arguments are provided in the script.

Conclusion

This lecture provided a short and incomplete introduction to Python.

Many Python libraries that will not be introduced but might be relevant in this lecture:

- **Pandas**: data management and analysis tools
- **Scikit-learn**: basic machine learning and data analysis tools
- **Matplotlib**: data plotting tools
- **Tensorflow** and **Keras**: deep-learning tools
- ...

Next lecture: introduction to the first scenario about **Activity Monitoring**